

Data Vortex — Round 1, Phase 2

Rebuilding the Social Engine · Analytical Core

SQL reasoning over the restored Dataset 01 — 12 queries, live outputs

Team Forge-X · Kunal Choudhary

1 · Schema, and why it is shaped this way

The Phase-1 output is two flat CSVs. Loading them verbatim and reaching for `text_content LIKE '%#Tag%'` in every query would make the string format part of the analytic logic, so three structural decisions come first. Full rationale is in `docs/schema_design.md`.

Structure	Reasoning it encodes
<code>posts.post_id</code> is a real PRIMARY KEY	it was NOT a key in the raw file (352 ids repeated across 712 rows). Making it one now means the dedup claim cannot silently regress.
<code>CHECK (platform IN (...))</code> on an enumerated set	the raw column held '' and 'NULL' as values; a bad future export now fails at <code>INSERT</code> instead of becoming a sixth platform.
<code>CHECK (likes >= 0)</code> with NULL permitted	the domain rule is in the schema, and NULL is deliberately allowed: a missing count is a known unknown, 0 is a claim nobody liked the post.
<code>posts.user_id → users.user_id</code> foreign key	referential integrity becomes a constraint, not an EDA task. <code>PRAGMA foreign_key_check</code> returns 0 rows.
<code>post_tags(post_id, tag)</code> as a long-form relation	hashtag analysis becomes <code>GROUP BY</code> , and the co-occurrence self-join in Q10 only exists because tags are rows.
location split into city + country	'Berlin, Germany' cannot be grouped without parsing, and one string hid UK vs United Kingdom.
<code>has_time</code> flag carried into SQL	date-only intake rows have no clock. Keeping the guard in the schema means an hour-of-day query has to opt in to the artefact rather than stumble into it.
<code>v_posts_enriched</code> view defines engagement once	SQL's + returns NULL if ANY operand is NULL, so <code>AVG(likes+shares+comments)</code> would quietly drop the 1,535 rows with missing likes. <code>COALESCE</code> is defined in exactly one place, in a view, so no query can forget it.

2 · Queries, outputs and the reasoning

Outputs below are captured by `src/run_sql.py` from `output/social_engine.db` at build time — the same tables appear as terminal screenshots and as `output/sql_outputs.md`. Nothing is transcribed, which is what the “hardcoded outputs will lead to disqualification” clause is testing.

Q1 · TREND DETECTION

monthly volume, momentum, and how anomalous each

```
● ● ● Q1 · TREND DETECTION SQLite 3.40.1

sqlite> .read queries/challenges.sql -- Q1
WITH monthly AS (
  SELECT * FROM v_monthly
),
stats AS (
  SELECT AVG(posts) AS mu_posts,
         SQRT(AVG(posts * posts) - AVG(posts) * AVG(posts)) AS sd_posts
  FROM monthly
)
SELECT m.post_month,
       m.posts,
       LAG(m.posts) OVER (ORDER BY m.post_month) AS prev_month,
       ROUND(100.0 * (m.posts - LAG(m.posts) OVER (ORDER BY m.post_month))
             / NULLIF(LAG(m.posts) OVER (ORDER BY m.post_month), 0), 2) AS mom_pct,
       ROUND(m.mean_engagement, 1) AS mean_engagement,
       ROUND((m.posts - s.mu_posts) / NULLIF(s.sd_posts, 0), 2) AS z_volume,
       RANK() OVER (ORDER BY m.posts DESC) AS volume_rank
FROM monthly m CROSS JOIN stats s
ORDER BY m.post_month
```

post_month	posts	prev_month	mom_pct	mean_engagement	z_volume	volume_rank
2024-05	878	NULL	NULL	3,561	1.07	1
2024-06	841	878	-4.21	3,528.2	-0.44	9
2024-07	862	841	2.5	3,545.8	0.42	7
2024-08	875	862	1.51	3,738.5	0.95	2
2024-09	822	875	-6.06	3,569.9	-1.21	11
2024-10	871	822	5.96	3,646.3	0.79	3
2024-11	865	871	-0.69	3,624.8	0.54	6
2024-12	867	865	0.23	3,686.9	0.62	5
2025-01	846	867	-2.42	3,617.5	-0.23	8
2025-02	794	846	-6.15	3,639.5	-2.36	12
2025-03	869	794	9.45	3,688.6	0.7	4
2025-04	831	869	-4.37	3,632.6	-0.85	10

```
rows: 12 · 12.61 ms · output/social_engine.db
```

Logic — `v_monthly` gives the grain (one row per month). `LAG` supplies the previous month for momentum. The series mean/stddev are computed by aggregate-over-window so every row is compared against the whole year without a second pass; stddev is written out manually because SQLite has no `STDDEV` -- [PG] `stddev_pop(total_engagement) OVER ()`.

Q2 · TREND DETECTION

3-month centred moving average and the inflection

Q2 · TREND DETECTION

SQLite 3.40.1

```
sqlite> .read queries/challenges.sql -- Q2
WITH ordered AS (
  SELECT post_month, posts,
    AVG(posts) OVER (
      ORDER BY post_month
      ROWS BETWEEN 1 PRECEDING AND 1 FOLLOWING) AS ma3,
    ROW_NUMBER() OVER (ORDER BY post_month) AS mo
  FROM v_monthly
),
turned AS (
  SELECT o.*,
    SIGN(ma3 - LAG(ma3) OVER (ORDER BY mo)) AS prev_dir,
    SIGN(LEAD(ma3) OVER (ORDER BY mo) - ma3) AS next_dir
  FROM ordered o
)
SELECT post_month, posts, ROUND(ma3, 1) AS ma3,
  CASE WHEN prev_dir = 1 AND next_dir = -1 THEN 'PEAK'
  WHEN prev_dir = -1 AND next_dir = 1 THEN 'TROUGH'
  ELSE '-' END AS inflection
FROM turned
ORDER BY post_month
```

post_month	posts	ma3	inflection
2024-05	878	859.5	--
2024-06	841	860.3	PEAK
2024-07	862	859.3	--
2024-08	875	853	TROUGH
2024-09	822	856	PEAK
2024-10	871	852.7	TROUGH
2024-11	865	867.7	PEAK
2024-12	867	859.3	--
2025-01	846	835.7	TROUGH
2025-02	794	836.3	PEAK
2025-03	869	831.3	TROUGH
2025-04	831	850	--

rows: 12 · 10.29 ms · output/social_engine.db

Logic — a frame of 1 PRECEDING / 1 FOLLOWING centres the average, which is the right choice for finding a turn (a trailing average lags the turn by half its window). DENSE_RANK over the calendar gives the ordinal month so FIRST_VALUE/LAST_VALUE can name the endpoints.

Q3 · BEHAVIOURAL GROUPING

platform league table with the intake gap

Q3 · BEHAVIOURAL GROUPING

SQLite 3.40.1

```
sqlite> .read queries/challenges.sql -- Q3
SELECT platform,
  COUNT(*) AS posts,
  ROUND(100.0 * COUNT(*) / SUM(COUNT(*) OVER ()), 2) AS share_pct,
  SUM(CASE WHEN likes IS NULL THEN 1 ELSE 0 END) AS likes_missing,
  ROUND(100.0 * SUM(CASE WHEN likes IS NULL THEN 1 ELSE 0 END) / COUNT(*), 2) AS likes_missing_pct,
  ROUND(AVG(engagement), 1) AS mean_engagement,
  ROUND(AVG(shares), 1) AS mean_shares,
  ROUND(AVG(comments), 1) AS mean_comments,
  CASE WHEN platform = 'Unspecified' THEN 0 ELSE 1 END AS in_rate_comparison
FROM v_posts_enriched
GROUP BY platform
ORDER BY posts DESC
```

platform	posts	share_pct	likes_missing	likes_missing_pct	mean_engagement	mean_shares	mean_comments	in_rate_comparison
YouTube	1770	17.32	275	15.54	3,621.5	1,009	505.3	1
Facebook	1763	17.25	269	15.26	3,636.1	986.8	507.6	1
Twitter	1742	17.04	273	15.67	3,561.5	1,006.8	503.5	1
Reddit	1716	16.79	243	14.16	3,606.9	985.1	506.9	1
Instagram	1689	16.52	254	15.04	3,673.2	1,043.3	501.8	1
Unspecified	1541	15.08	218	14.15	3,646.9	999.2	500.4	0

rows: 6 · 10.65 ms · output/social_engine.db

Logic — every platform gets its own means, and in_rate_comparison marks which rows are admissible for cross-platform claims, instead of quietly NULLing a column (which a reader would mistake for missing data). Rates are per-post means over that platform's own rows, so the 15.08% intake gap cannot distort a total and a rate in the same table.

Q4 · ANOMALY DISCOVERY

amplification anomalies -- posts that are shared

```
●●● Q4 · ANOMALY DISCOVERY SQLite 3.40.1

sqlite> .read queries/challenges.sql -- Q4
WITH flagged AS (
  SELECT post_id, user_id, platform, post_month, likes, shares, comments,
         amplify_ratio
  FROM v_posts_enriched
  WHERE likes > 0 AND amplify_ratio > 3
)
SELECT COUNT(*) AS anomalous_posts,
       (SELECT COUNT(*) FROM v_posts_enriched
        WHERE likes > 0) AS comparable_posts,
       ROUND(100.0 * COUNT(*) /
             (SELECT COUNT(*) FROM v_posts_enriched WHERE likes > 0), 2) AS pct_of_comparable,
       ROUND(AVG(amplify_ratio), 2) AS mean_ratio,
       ROUND(MAX(amplify_ratio), 1) AS worst_ratio
FROM flagged

anomalous_posts | comparable_posts | pct_of_comparable | mean_ratio | worst_ratio
846             | 8688             | 9.74              | 24.8       | 1,863

rows: 1 · 6.52 ms · output/social_engine.db
```

Logic — ratio-based anomaly, not magnitude-based, because every count here is uniform on [0,5000] so no row is an outlier by value. A row only becomes suspicious when the RELATIONSHIP between its columns breaks. likes > 0 is required or the ratio is undefined. 3x is the threshold stated up front so the query is reproducible, not tuned to a result.

Q5 · ANOMALY DISCOVERY

attribution -- is the amplification anomaly a

```
●●● Q5 · ANOMALY DISCOVERY SQLite 3.40.1

sqlite> .read queries/challenges.sql -- Q5
WITH per_user AS (
  SELECT u.user_id,
         COALESCE(a.anom, 0) AS anom
  FROM users u
  LEFT JOIN (SELECT user_id, COUNT(*) AS anom
             FROM v_posts_enriched
             WHERE likes > 0 AND amplify_ratio > 3
             GROUP BY user_id) a ON a.user_id = u.user_id
),
params AS (
  SELECT COUNT(*) AS n_users,
         SUM(anom) AS n_anom,
         SUM(anom) * 1.0 / COUNT(*) AS lambda_
  FROM per_user
),
observed AS (
  SELECT SUM(anom >= 3) AS obs_ge3,
         SUM(anom >= 4) AS obs_ge4,
         MAX(anom) AS obs_max
  FROM per_user
),
expected AS (
  SELECT n_users, lambda_,
         n_users * EXP(-1.0 * lambda_) * (1.0 + lambda_
         + POW(lambda_, 2) / 2.0) AS cum_le2,
         n_users * EXP(-1.0 * lambda_) * (1.0 + lambda_
         + POW(lambda_, 2) / 2.0 + POW(lambda_, 3) / 6.0) AS cum_le3
  FROM params
)
SELECT p.n_users,
       p.n_anom,
       ROUND(p.lambda_, 4) AS poisson_lambda,
       o.obs_ge3,
       ROUND(p.n_users - e.cum_le2, 1) AS expected_ge3_by_chance,
       o.obs_ge4,
       ROUND(p.n_users - e.cum_le3, 1) AS expected_ge4_by_chance,
       o.obs_max AS busiest_author,
       ROUND(o.obs_ge3 * 1.0 / NULLIF(p.n_users - e.cum_le2, 0), 3) AS observed_over_expected,
       CASE WHEN o.obs_ge3 BETWEEN 0.8 * (p.n_users - e.cum_le2)
              AND 1.25 * (p.n_users - e.cum_le2)
            THEN 'DIFFUSE -- indistinguishable from random scatter, NOT a coordinated ring'
            ELSE 'CONCENTRATED -- attribution is warranted' END AS verdict
FROM params p CROSS JOIN observed o CROSS JOIN expected e

n_users | n_anom | poisson_lambda | obs_ge3 | expected_ge3_by_chance | obs_ge4 | expected_ge4_by_chance | busiest_author | observed_over_expected | verdict
1500    | 846    | 0.56           | 31      | 29.6                   | 2        | 4                      | 4              | 1.05                   | DIFFUSE -- indistinguishable from random scatter, NOT a coordinated ring

rows: 1 · 8.31 ms · output/social_engine.db
```

Logic — a per-user ranking would "find" a ring in ANY diffuse anomaly, so instead we test the observed tail against a Poisson null: if 846 anomalies were scattered at random over 1500 authors, the count per author is Poisson($\lambda = n_anomalies / n_users$), and $P(X \geq k)$ has a closed form. Anything within a few percent of that expectation is not evidence of a ring. This is the query that stops a plausible-looking "bot network" finding from being published.

Q6 · DATA-QUALITY FORENSICS

the midnight illusion -- all 24 hours

```
● ● ● Q6 · DATA-QUALITY FORENSICS SQLite 3.40.1

sqlite> .read queries/challenges.sql -- Q6
SELECT h.post_hour,
       COUNT(*) AS rows_naive,
       ROUND(100.0 * COUNT(*) / SUM(COUNT(*) OVER ()), 2) AS pct_naive,
       SUM(h.has_time) AS rows_with_real_time,
       ROUND(100.0 * SUM(h.has_time) / SUM(SUM(h.has_time) OVER ()), 2) AS pct_valid,
       COUNT(*) - SUM(h.has_time) AS artefact_rows
FROM v_posts_enriched h
GROUP BY h.post_hour
ORDER BY artefact_rows DESC, h.post_hour
```

post_hour	rows_naive	pct_naive	rows_with_real_time	pct_valid	artefact_rows
0	3325	32.53	323	4.47	3002
1	311	3.04	311	4.31	0
2	301	2.94	301	4.17	0
3	318	3.11	318	4.41	0
4	286	2.8	286	3.96	0
5	279	2.73	279	3.86	0
6	295	2.89	295	4.09	0
7	281	2.75	281	3.89	0
8	294	2.88	294	4.07	0
9	307	3	307	4.25	0
10	305	2.98	305	4.22	0
11	292	2.86	292	4.04	0

```
rows: 24 (showing first 12 of 24) · 3.28 ms · output/social_engine.db
```

Logic — the naive distribution is counted from every row; the valid one from has_time = 1. The difference attributable to the format is named explicitly, which is what turns "the data looks like X" into "the data looks like X because the intake serialised Y".

Q7 · BEHAVIOURAL GROUPING

RFM segmentation of the author base

```
● ● ● Q7 · BEHAVIOURAL GROUPING SQLite 3.40.1

sqlite> .read queries/challenges.sql -- Q7
WITH base AS (
  SELECT user_id,
         CAST(JULIANDAY((SELECT MAX(posted_at) FROM posts))
              - JULIANDAY(MAX(posted_at))) AS INTEGER) AS recency_days,
         COUNT(*) AS frequency,
         SUM(engagement) AS monetary
  FROM v_posts_enriched
  GROUP BY user_id
),
scored AS (
  SELECT b.*,
         NTILE(5) OVER (ORDER BY recency_days DESC) AS r_tile, -- small days = recent
         NTILE(5) OVER (ORDER BY frequency) AS f_tile,
         NTILE(5) OVER (ORDER BY monetary) AS m_tile
  FROM base b
)
SELECT CASE
  WHEN r_tile >= 4 AND f_tile >= 4 AND m_tile >= 4 THEN 'champions'
  WHEN r_tile >= 4 AND f_tile <= 2 THEN 'new_and_promising'
  WHEN r_tile <= 2 AND f_tile >= 4 THEN 'drifting_regulars'
  WHEN r_tile <= 2 AND m_tile <= 2 THEN 'lapsed_low_value'
  ELSE 'mid_value'
END AS segment,
COUNT(*) AS users,
ROUND(100.0 * COUNT(*) / SUM(COUNT(*)) OVER (), 2) AS share_pct,
ROUND(AVG(frequency), 2) AS avg_posts,
ROUND(AVG(monetary), 1) AS avg_total_engagement,
ROUND(AVG(recency_days), 0) AS avg_days_since_last_post
FROM scored
GROUP BY segment
ORDER BY users DESC

segment | users | share_pct | avg_posts | avg_total_engagement | avg_days_since_last_post
mid_value | 592 | 39.47 | 6.94 | 25,416.4 | 45
lapsed_low_value | 310 | 20.67 | 4.13 | 13,702.1 | 117
champions | 281 | 18.73 | 9.72 | 36,808.7 | 12
new_and_promising | 167 | 11.13 | 4.53 | 15,191.6 | 12
drifting_regulars | 150 | 10 | 8.95 | 32,423.8 | 82

rows: 5 · 16.44 ms · output/social_engine.db
```

Logic — Recency/Frequency/Monetary are scored with NTILE(5), which buckets by rank so a few extreme users cannot drag the cut points around the way fixed thresholds would. Concatenating the three digits gives a human-readable segment label; SUM OVER supplies each segment's share.

Q8 · CORRELATION ANALYSIS

does follower count buy engagement?

Q8 · CORRELATION ANALYSIS

SQLite 3.40.1

```
sqlite> .read queries/challenges.sql -- Q8
WITH per_user AS (
  SELECT p.user_id,
         AVG(p.engagement) AS mean_engagement,
         u.follower_count
  FROM v_posts_enriched p
  JOIN users u ON u.user_id = p.user_id
  GROUP BY p.user_id, u.follower_count
),
agg AS (
  SELECT COUNT(*) AS n,
         SUM(follower_count) AS sx,
         SUM(mean_engagement) AS sy,
         SUM(follower_count * mean_engagement) AS sxy,
         SUM(follower_count * follower_count) AS sxx,
         SUM(mean_engagement * mean_engagement) AS syy
  FROM per_user
)
SELECT n AS users,
       ROUND((n*sxy - sx*sy) /
              (SQRT((n*sxx - sx*sx) * (n*syy - sy*sy))), 4) AS pearson_r,
       ROUND(((n*sxy - sx*sy) / (SQRT((n*sxx - sx*sx) * (n*syy - sy*sy))))
              * ((n*sxy - sx*sy) / (SQRT((n*sxx - sx*sx) * (n*syy - sy*sy))))
              * 100, 2) AS r_squared_pct,
       CASE WHEN ABS((n*sxy - sx*sy) /
                     (SQRT((n*sxx - sx*sx) * (n*syy - sy*sy)))) < 0.1
              THEN 'no material relationship'
              ELSE 'relationship worth modelling' END AS verdict
FROM agg
```

users	pearson_r	r_squared_pct	verdict
1500	-0.04	0.19	no material relationship

rows: 1 · 13.78 ms · output/social_engine.db

Logic — SQLite has no CORR(), so the Pearson product-moment coefficient is assembled from raw sums -- the definition, not a shortcut. Dividing by $n*(n-1)$ style denominators would be wrong; the covariance term ($n*S_{xy} - S_x*S_y$) and the two variance terms must share the same n . [PG] CORR(follower_count, mean_engagement)

Q9 · CORRELATION ANALYSIS

same question, bucketed -- because a near-zero

Q9 · CORRELATION ANALYSIS

SQLite 3.40.1

```
sqlite> .read queries/challenges.sql -- Q9
WITH ranked AS (
  SELECT u.follower_count,
         NTILE(5) OVER (ORDER BY u.follower_count) AS quintile,
         p.engagement
  FROM users u JOIN v_posts_enriched p ON p.user_id = u.user_id
)
SELECT quintile,
       CASE quintile WHEN 1 THEN '1 (fewest)' WHEN 2 THEN '2' WHEN 3 THEN '3'
                    WHEN 4 THEN '4' ELSE '5 (most)' END AS follower_band,
       MIN(follower_count) AS min_followers,
       MAX(follower_count) AS max_followers,
       COUNT(*) AS posts,
       ROUND(AVG(engagement), 1) AS mean_engagement
FROM ranked
GROUP BY quintile
ORDER BY quintile
```

quintile	follower_band	min_followers	max_followers	posts	mean_engagement
1	1 (fewest)	109	10486	2045	3,630.8
2	2	10486	19785	2044	3,684.7
3	3	19785	29790	2044	3,607
4	4	29790	39885	2044	3,621.2
5	5 (most)	39894	49944	2044	3,574.8

rows: 5 · 19.94 ms · output/social_engine.db

Logic — follower quintiles via NTILE, then mean engagement per bucket. If the relationship were monotone the buckets would be ordered; a flat band is the evidence that $r=0$ is a real null, not a modelling error.

Q10 · TREND + GROUPING

hashtag momentum -- which tag accelerated fastest

Q10 · TREND + GROUPING

SQLite 3.40.1

```
sqlite> .read queries/challenges.sql -- Q10
WITH tag_month AS (
  SELECT pt.tag, p.post_month,
         COUNT(DISTINCT p.post_id) AS posts
  FROM post_tags pt JOIN posts p ON p.post_id = pt.post_id
  GROUP BY pt.tag, p.post_month
),
share AS (
  SELECT tm.*,
         ROUND(100.0 * tm.posts / m.month_posts, 2) AS pct_of_month,
         m.month_posts
  FROM tag_month tm
  JOIN (SELECT post_month, COUNT(*) AS month_posts FROM posts GROUP BY post_month) m
       ON m.post_month = tm.post_month
),
pivoted AS (
  SELECT s.*,
         LAG(pct_of_month) OVER (PARTITION BY tag ORDER BY post_month) AS prev_pct
  FROM share s
)
SELECT tag,
       MAX(CASE WHEN post_month = (SELECT MAX(post_month) FROM posts) THEN pct_of_month END) AS pct_final_month,
       MAX(CASE WHEN post_month = (SELECT MIN(post_month) FROM posts) THEN pct_of_month END) AS pct_first_month,
       ROUND(MAX(CASE WHEN post_month = (SELECT MAX(post_month) FROM posts) THEN pct_of_month END)
            - MAX(CASE WHEN post_month = (SELECT MIN(post_month) FROM posts) THEN pct_of_month END), 2)
         AS pp_change,
       MAX(pct_of_month - prev_pct) AS best_single_month_gain_pp
FROM pivoted
GROUP BY tag
ORDER BY pp_change DESC
LIMIT 8
```

tag	pct_final_month	pct_first_month	pp_change	best_single_month_gain_p
Tech	8.18	5.47	2.71	2.89
Sustainable	8.9	6.38	2.52	1.71
ProductLaunch	8.06	5.81	2.25	2.64
TrendAlert	7.46	6.04	1.42	1.86
Travel	7.82	6.61	1.21	1.61
Promo	7.7	6.72	0.98	2
Fitness	7.7	6.83	0.87	1.73
Limited	6.62	5.92	0.7	1.81

rows: 8 · 47.90 ms · output/social_engine.db

Logic — share-of-month is the correct base rate, otherwise a tag only looks like it is growing because the whole corpus grew. The self-join on the prior month replaces a LAG that cannot be used directly here, because months where a tag has zero posts have no row to lag.

Q11 · ANOMALY DISCOVERY (gap & islands)

posting streaks per user

```
● ● ● Q11 · ANOMALY DISCOVERY (gap & islands) SQLite 3.40.1

sqlite> .read queries/challenges.sql -- Q11
WITH days AS (
  SELECT user_id, post_date,
         CAST(JULIANDAY(post_date) AS INTEGER) AS d
  FROM posts
  GROUP BY user_id, post_date
),
grp AS (
  SELECT user_id, d,
         d - ROW_NUMBER() OVER (PARTITION BY user_id ORDER BY d) AS island
  FROM days
)
SELECT streak_len,
       COUNT(*) AS streaks,
       ROUND(AVG(posts_in_streak),1) AS avg_posts_per_day,
       MAX(posts_in_streak) AS max_posts_in_one_day
FROM (
  SELECT user_id, island,
         COUNT(*) AS posts_in_streak,
         COUNT(*) AS streak_len
  FROM grp GROUP BY user_id, island
)
GROUP BY streak_len
ORDER BY streak_len

streak_len | streaks | avg_posts_per_day | max_posts_in_one_day
1          | 9763    | 1                 | 1
2          | 179     | 2                 | 2
3          | 4       | 3                 | 3

rows: 3 · 33.02 ms · output/social_engine.db
```

Logic — classic islands technique. ROW_NUMBER over consecutive calendar dates and subtracting the day offset makes the difference CONSTANT inside a run, so GROUP BY that expression collapses each streak. A streak is a behaviour, not a value -- exactly the shape SQL is good at.

Q12 · INTEGRITY GATE

prove the restored tables can answer anything above

```
● ● ● Q12 · INTEGRITY GATE SQLite 3.40.1

sqlite> .read queries/challenges.sql -- Q12
SELECT 'duplicate post id' AS check_name,
       COUNT(*) - COUNT(DISTINCT post_id) AS violations FROM posts
UNION ALL
SELECT 'orphan user id',
       (SELECT COUNT(*) FROM posts p
        WHERE NOT EXISTS (SELECT 1 FROM users u WHERE u.user_id = p.user_id))
UNION ALL
SELECT 'negative engagement',
       (SELECT COUNT(*) FROM posts
        WHERE COALESCE(likes,0) < 0 OR COALESCE(shares,0) < 0
              OR COALESCE(comments,0) < 0)
UNION ALL
SELECT 'post before signup',
       (SELECT COUNT(*) FROM posts p JOIN users u USING (user_id)
        WHERE p.posted_at < u.account_created)
UNION ALL
SELECT 'timestamps outside window',
       (SELECT COUNT(*) FROM posts
        WHERE posted_at < '2024-05-01' OR posted_at >= '2025-05-01')

check_name | violations
duplicate post id | 0
orphan user id | 0
negative engagement | 0
post before signup | 0
timestamps outside windo | 0

rows: 5 · 14.34 ms · output/social_engine.db
```

Logic — five assertions that would each invalidate a section of this file. Returning 0 for violations is the proof, so the analytic work is auditable rather than trusted. A judge can run this one query alone.

3 · What the SQL phase establishes

- Monthly volume mean 852, population SD 24.5. z-scores run -2.36 (2025-02) to +1.07 (2024-05), and 3 of 12 months sit beyond +/-1 SD. The rest oscillate around a flat mean with no monotone component, so neither a growth narrative nor a 'surge' month is supported. The window function computes the year's own mean and SD per row, so the comparison needs no second scan and the verdict is a property of the data rather than of a chosen threshold.
- A centred 3-month frame (1 PRECEDING/1 FOLLOWING) is used deliberately: a trailing average lags a turning point by half its window and would name the wrong month as the peak. Q1 and Q2 disagree by exactly that much, and Q2 is the defensible one.
- Per-1000-posts rates and total volume come from one pass. Reporting a rate over 85% of rows against a total over 100% is the classic way to manufacture a platform difference; keeping 'Unspecified' visible prevents it.
- Every count here is uniform on [0,5000], so no row is an outlier by MAGNITUDE. The anomaly only exists in the RELATIONSHIP between columns, which is why the ratio — not a percentile on likes — is the correct detector.
- This is the query that protects the submission. A per-user ranking would 'discover' a bot ring in any diffuse anomaly, so the observed tail is tested against a Poisson null with $\lambda = n_{\text{anom}}/n_{\text{users}}$. Observed-to-expected ≈ 1.0 means the amplification anomaly is spread at random: it is an artefact of independently generated columns, NOT a coordinated ring. The finding is reported as a negative result.
- The 00:00 row is the whole argument of this report in one line: 3,325 rows naively, 3,002 of them pure format artefact, leaving a flat real distribution. An analyst who does not know the intake formats cannot tell these apart.
- RFM over 1,500 authors splits them as mid_value 592, lapsed_low_value 310, champions 281, new_and_promising 167, drifting_regulars 150. NTILE(5) buckets by rank so a few extreme authors cannot drag cut points around the way fixed thresholds would, and SUM(COUNT(*)) OVER () gives each segment's share in the same pass. The segments separate volume from yield, which is the distinction the platform could actually act on.
- SQLite has no CORR(), so Pearson r is assembled from raw cross-products — the definition, not a shortcut. $r = -0.0433$ ($r^2 \approx 0.19\%$) is a genuine null: audience size does not predict resonance in this corpus.
- A near-zero r can hide a non-monotone shape, so the same relation is re-tested in follower quintiles. Mean engagement moves 3,574.8 to 3,684.7 across five bands with no ordering — flat. That confirms Q8's null instead of trusting one coefficient.
- Share-of-month is the correct base rate; raw tag counts would report the corpus's own growth as tag momentum. LAG cannot be used directly because a month with zero posts for a tag has no row to lag, hence the pivot with MAX(CASE...) — the standard sparse-series fix.
- Gap-and-islands via ROW_NUMBER over calendar days: subtracting the row number makes the offset constant inside a run, so GROUP BY collapses each streak. Streaking is a behaviour, not a value, and this is the shape SQL handles that spreadsheets do not.
- Five assertions that would each invalidate a section above — duplicate keys, orphans, negative engagement, time-travelling posts, out-of-window dates — all return 0. The analytics are auditable rather than trusted.

4 · Cross-cutting conclusion

Taken together the twelve queries describe a corpus that is **uniform by construction**: platform, sentiment, geography, audience size and hashtag choice all explain close to nothing about engagement, and the only large apparent signals — a midnight posting culture, an amplification ring, growth momentum — each dissolve under a data-quality or statistical-control check. For a restored pipeline that is the correct result to report. The defects were in the transport layer (formats, entities, sign bits, duplicate replays), not in the sampling, so the analysis is sound precisely because nothing significant survives the corrections.

Reproduce it

```
python src/clean_data.py && python src/build_db.py && python src/run_sql.py && python src/make_screenshots.py  
&& python src/build_report.py
```